

Hadoop Virtual Machine Guide

CSE 4/587

Prepared by Jacob Condello

VM SETUP

1. Download and install VirtualBox [1]
2. Download the virtual machine from UBbox [2]
3. Import Hadoop.ova by either double clicking it or from VirtualBox selecting File->Import Appliance
4. After setup run the virtual machine named "Hadoop"
5. Login with user "hadoop" password "hadoop"

HADOOP SETUP

1. Make sure you have access to a text editor. The VM comes with mousepad or you can install a different one.
2. Add the following lines to the end of your ~/.bashrc file:

```
export JAVA_HOME=/usr/lib/jvm/java-6-openjdk-i386
export HADOOP_CLASSPATH=$JAVA_HOME/lib/tools.jar
```
3. Apply changes by running:

```
source ~/.bashrc
```
4. Start Hadoop by running:

```
start-hadoop.sh
```
5. Make a new directory on hdfs for the input files:

```
hdfs dfs -mkdir -p ~/input/
```
6. Copy the local directory books/ to the hdfs directory input/:

```
hdfs dfs -put ~/books/ ~/input
```
7. Make sure you copied them to the right place:

```
hdfs dfs -ls ~/input
```

MR CODE

1. Create a file named WordCount.java and copy in the source code for Example: WordCount v1.0 from hadoop.achache.org [3]
2. Run the following to compile WordCount.java:

```
hadoop com.sun.tools.javac.Main WordCount.java
```

Note: If you get an error at this point, it is likely either your java version, JAVA_HOME, or HADOOP_CLASSPATH that aren't configured correctly.

3. Run the following to create a jar:

```
jar cf wc.jar WordCount*.class
```

4. You can now run your MR job with the following:

```
hadoop jar wc.jar WordCount ~/input/books ~/output1
```

Notice that each time you run a MR job you need to either specify a new output file or delete the old one

5. You can now view the results by running:

```
hdfs dfs -cat ~/output1/*
```

Or if you need the results locally, copy from hdfs then cat:

```
hdfs dfs -get ~/output1  
cat part-r-00000
```

REFERENCES

[1] <https://www.virtualbox.org/wiki/Downloads>

[2] <https://buffalo.app.box.com/s/t72vy6py6r6v852anwa55whl4mmibfaw>

[3] <https://hadoop.apache.org/docs/stable/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>